

Algorithms Crash Review

Laura Toma, csci2200, Bowdoin College

Here is a list of topics studied in algorithms this semester:

1. Worst case, best case running times.
 - e.g. Binary search runs in logarithmic time worst case and $O(1)$ best case.
2. Asymptotic growth of functions (O, Ω, Θ). Used to express efficiency (running time, space).
 - This algorithm runs in $O(n^2)$ time and uses $\Theta(n)$ space.
 - This algorithm runs in $\Theta(n \lg n)$ worst case.
3. Recurrences. Used to analyze complexity of recursive algorithms.
 - Example: mergesort recurrence: $T(n) = 2T(n/2) + \Theta(n)$ solves to $\Theta(n \lg n)$.
4. Comparison-based sorting
 - Quadratic sorts: insertion sort, bubble sort, selection sort
 - Mergesort, Heapsort, Quicksort, Randomized quicksort
 - Comparison-based sorting lower bound: Any sorting algorithm that uses only comparisons to order the elements must take $\Omega(n \lg n)$ in the worst case.
5. Sorting without comparisons
 - Not a general purpose sort; makes additional assumptions, usually that the keys to be sorted are integers in a small range.
 - Example: sort n keys that are all integers in the range $\{-10, \dots, 10\}$.
 - Counting sort, Radix sort, Bucket sort
 - Runs in $O(n + k)$ where k is the range
6. Selection: given a set of n elements, find the i th smallest.
 - Idea: partition and recurse. Quick-Select in expected $O(n)$ time ; $O(n)$ Smart-Select in worst-case $O(n)$ time. Theoretically better but in practice Quick-Select is preferred.
7. The binary heap
 - Implements the priority queue ADT, and supports FindMin (in $O(1)$ time), DeleteMin, Insert, ChangeKey all in $O(\lg n)$ time
 - A max-heap is defined symmetrically

8. Binary search trees (BST) and red-black trees

- A BST supports insert, delete, search, min, max, pred, succ, all in $O(h)$ time; tree walks (in-order, post-order, pre-order) in $O(n)$ time.
- A red-black tree is a BST plus additional invariants that ensure that $h = \Theta(\lg n)$.
- On red-black-trees all above operations run in $O(\lg n)$ time which includes the additional time to maintain $h = \Theta(\lg n)$ via rotations

9. Technique: Divide-and-conquer

- Technique that solves a problem by decomposing it into smaller subproblems and solving each one recursively, and then “merging” their solutions
- E.g. mergesort, maximum partial sum, counting inversions.
- Also: binary search, finding missing number

10. Technique Dynamic programming

- Technique used for solving optimization problems that have optimal substructure and overlapping subproblems. Store (cache) solutions to subproblems in a table to avoid recomputation.
- Examples: rod cutting, house robber, 0-1 knapsack, longest common subsequence, weighted interval scheduling, subset sum, weighted subset sum (0-1 knapsack), longest common subsequence; skis and skiers, also shortest paths (Bellman Ford)

11. Technique: Greedy algorithms

- Instead of dynamic programming (which does an exhaustive search), make a choice “locally” and recurse on what’s left. The choice is determined by examining only local information, without going into recursion to see how good the choice is “globally”.
- Examples: fractional knapsack, activity selection, art gallery guarding. Also Dijkstra SSSP, Prim and Kruskal’s algorithm for MST are all greedy.
- Greedy heuristics are used a lot in AI but not so much in algorithms (they rarely guarantee the optimal solution)

12. Graph algorithms

- Graph representation (adjacency list, adjacency matrix)
- Graph traversal: BFS and DFS
 - G can be directed or undirected
 - run in linear time $O(V + E)$
 - BFS used to: find connected components, reachability, check bipartiteness (G undirected), compute shortest paths (G un-weighted, all edges have weight 1)
 - DFS used to: find connected components, reachability, find directed cycles, topological sort (G must be a DAG)
- Directed acyclic graphs (DAGs)

- Topological order: Order of the vertices such that for any edge (x, y) , vertex x comes before vertex y in topological order (all edges are “forward”)
- Any DAG can be topologically ordered in linear $O(V + E)$ time either directly or via DFS
- Many problems have easier/faster solutions on DAGs. E.g. SSSP on DAGs can be computed in $O(V + E)$ time; Longest paths on DAGs in $O(V + E)$ (note: on general graphs longest path is an NPC problem).
- Shortest paths (SSSP): Find shortest paths from a vertex to all other vertices
 - On DAGs: $O(V + E)$ time
 - On general graphs with non-negative weights: Dijkstra’s runs in $O(E \lg V)$ time
 - On general digraphs without negative cycles: Bellman-Ford’s algorithm runs in $O(V \cdot E)$ and can also detect negative-weight cycles
- Minimum spanning tree (MST)
 - G : connected, undirected, weighted
 - MST can be computed in $O(E \lg V)$ time with Prim’s or Kruskal’s algorithms
 - Union-find data structure: supports Find and Join/Union operations.

13. Complexity:

- P : problems that can be solved in polynomial time (on a deterministic Turing machine)
- NP : problems that can be verified in polynomial time (equivalent: solved in polynomial time on a non-deterministic Turing machine).
- NPC : a problem is in NPC if it is in NP and all problems in NP reduce to it in polynomial time. Intuitively, NPC is the core of hard problems in NP.
- Some NPC problems: SAT (satisfiability: is there an assignment that makes a given formula true), traveling salesman (TSP: find minimum-cost tour), longest path (find the longest simple path in a graph); find the largest clique subgraph (a clique is a complete graph).
- All NPC problems reduce to each other therefore if any of them is shown to be in P, it follows they all are and $P = NP$
- The \$1M questions: Is $P = NP$? Not known. Most people believe the answer is NO.